

# Contents

<b>1 BPF Instruction Set Specification, v1.0</b>	<b>1</b>
1.1 Documentation conventions	1
1.1.1 Types	1
1.1.2 Functions	2
1.1.3 Definitions	2
1.2 Instruction encoding	2
1.2.1 Instruction classes	3
1.3 Arithmetic and jump instructions	4
1.3.1 Arithmetic instructions	4
1.3.2 Byte swap instructions	6
1.3.3 Jump instructions	7
1.3.3.1 Helper functions	8
1.3.3.2 Program-local functions	8
1.4 Load and store instructions	9
1.4.1 Regular load and store operations	9
1.4.2 Sign-extension load operations	10
1.4.3 Atomic operations	10
1.4.4 64-bit immediate instructions	11
1.4.4.1 Maps	11
1.4.4.2 Platform Variables	12
1.4.5 Legacy BPF Packet access instructions	12

## 1 BPF Instruction Set Specification, v1.0

This document specifies version 1.0 of the BPF instruction set.

### 1.1 Documentation conventions

For brevity and consistency, this document refers to families of types using a shorthand syntax and refers to several expository, mnemonic functions when describing the semantics of instructions. The range of valid values for those types and the semantics of those functions are defined in the following subsections.

#### 1.1.1 Types

This document refers to integer types with the notation  $SN$  to specify a type's signedness ( $S$ ) and bit width ( $N$ ), respectively.

Meaning of signedness notation.

$S$	Meaning
$u$	unsigned
$s$	signed

Meaning of bit-width notation.

$N$	Bit width
-----	-----------

8	8 bits
16	16 bits
32	32 bits
64	64 bits
128	128 bits

For example, *u32* is a type whose valid values are all the 32-bit unsigned numbers and *s16* is a types whose valid values are all the 16-bit signed numbers.

### 1.1.2 Functions

- *htobe16*: Takes an unsigned 16-bit number in host-endian format and returns the equivalent number as an unsigned 16-bit number in big-endian format.
- *htobe32*: Takes an unsigned 32-bit number in host-endian format and returns the equivalent number as an unsigned 32-bit number in big-endian format.
- *htobe64*: Takes an unsigned 64-bit number in host-endian format and returns the equivalent number as an unsigned 64-bit number in big-endian format.
- *htole16*: Takes an unsigned 16-bit number in host-endian format and returns the equivalent number as an unsigned 16-bit number in little-endian format.
- *htole32*: Takes an unsigned 32-bit number in host-endian format and returns the equivalent number as an unsigned 32-bit number in little-endian format.
- *htole64*: Takes an unsigned 64-bit number in host-endian format and returns the equivalent number as an unsigned 64-bit number in little-endian format.
- *bswap16*: Takes an unsigned 16-bit number in either big- or little-endian format and returns the equivalent number with the same bit width but opposite endianness.
- *bswap32*: Takes an unsigned 32-bit number in either big- or little-endian format and returns the equivalent number with the same bit width but opposite endianness.
- *bswap64*: Takes an unsigned 64-bit number in either big- or little-endian format and returns the equivalent number with the same bit width but opposite endianness.

### 1.1.3 Definitions

#### Example

Sign extend an 8-bit number *A* to a 16-bit number *B* on a big-endian platform:

```
A:      10000110
B: 11111111 10000110
```

## 1.2 Instruction encoding

BPF has two instruction encodings:

- the basic instruction encoding, which uses 64 bits to encode an instruction
- the wide instruction encoding, which appends a second 64-bit immediate (i.e., constant) value after the basic instruction for a total of 128 bits.

The fields conforming an encoded basic instruction are stored in the following order:



BPF_LD	0x0 0	non-standard load operations	<a href="#">Load and store instructions</a>
BPF_LD X	0x0 1	load into register operations	<a href="#">Load and store instructions</a>
BPF_ST	0x0 2	store from immediate operations	<a href="#">Load and store instructions</a>
BPF_ST X	0x0 3	store from register operations	<a href="#">Load and store instructions</a>
BPF_AL U	0x0 4	32-bit arithmetic operations	<a href="#">Arithmetic and jump instructions</a>
BPF_JM P	0x0 5	64-bit jump operations	<a href="#">Arithmetic and jump instructions</a>
BPF_JM P32	0x0 6	32-bit jump operations	<a href="#">Arithmetic and jump instructions</a>
BPF_AL U64	0x0 7	64-bit arithmetic operations	<a href="#">Arithmetic and jump instructions</a>

### 1.3 Arithmetic and jump instructions

For arithmetic and jump instructions (BPF\_ALU, BPF\_ALU64, BPF\_JMP and BPF\_JMP32), the 8-bit 'opcode' field is divided into three parts:

4 bits (MSB)	1 bit	3 bits (LSB)
code	source	instruction class

#### code

the operation code, whose meaning varies by instruction class

#### source

the source operand location, which unless otherwise specified is one of:

source	value	description
BPF_K	0x00	use 32-bit 'imm' value as source operand
BPF_X	0x08	use 'src_reg' register value as source operand

#### instruction class

the instruction class (see [Instruction classes](#))

#### 1.3.1 Arithmetic instructions

BPF\_ALU uses 32-bit wide operands while BPF\_ALU64 uses 64-bit wide operands for otherwise identical operations. The 'code' field encodes the operation as below, where 'src' and 'dst' refer to the values of the source and destination registers, respectively.

code	value	offset	description
BPF_ADD	0x0 0	0	dst += src
BPF_SUB	0x1 0	0	dst -= src
BPF_MUL	0x2 0	0	dst *= src

BPF_DIV	0x3 0	0	$\text{dst} = (\text{src} \neq 0) ? (\text{dst} / \text{src}) : 0$
BPF_SDIV	0x3 0	1	$\text{dst} = (\text{src} \neq 0) ? (\text{dst} \text{ s} / \text{src}) : 0$
BPF_OR	0x4 0	0	$\text{dst}  = \text{src}$
BPF_AND	0x5 0	0	$\text{dst} \&= \text{src}$
BPF_LSH	0x6 0	0	$\text{dst} \ll= (\text{src} \& \text{mask})$
BPF_RSH	0x7 0	0	$\text{dst} \gg= (\text{src} \& \text{mask})$
BPF_NEG	0x8 0	0	$\text{dst} = -\text{dst}$
BPF_MOD	0x9 0	0	$\text{dst} = (\text{src} \neq 0) ? (\text{dst} \% \text{src}) : \text{dst}$
BPF_SMOD	0x9 0	1	$\text{dst} = (\text{src} \neq 0) ? (\text{dst} \text{ s}\% \text{src}) : \text{dst}$
BPF_XOR	0xa 0	0	$\text{dst} \wedge= \text{src}$
BPF_MOV	0xb 0	0	$\text{dst} = \text{src}$
BPF_MOVSX	0xb 0	8/16/32	$\text{dst} = (\text{s8}, \text{s16}, \text{s32})\text{src}$
BPF_ARSH	0xc 0	0	$\text{dst} \gg= (\text{src} \& \text{mask})$ <small><a href="#">term: sign extending&lt;Sign Extend&gt;</a></small>
BPF_END	0xd 0	0	byte swap operations (see <a href="#">Byte swap instructions</a> below)

Underflow and overflow are allowed during arithmetic operations, meaning the 64-bit or 32-bit value will wrap. If BPF program execution would result in division by zero, the destination register is instead set to zero. If execution would result in modulo by zero, for BPF\_ALU64 the value of the destination register is unchanged whereas for BPF\_ALU the upper 32 bits of the destination register are zeroed.

BPF\_ADD | BPF\_X | BPF\_ALU means:

```
dst = (u32) ((u32) dst + (u32) src)
```

where '(u32)' indicates that the upper 32 bits are zeroed.

BPF\_ADD | BPF\_X | BPF\_ALU64 means:

```
dst = dst + src
```

BPF\_XOR | BPF\_K | BPF\_ALU means:

```
dst = (u32) dst ^ (u32) imm32
```

BPF\_XOR | BPF\_K | BPF\_ALU64 means:

```
dst = dst ^ imm32
```

Note that most instructions have instruction offset of 0. Only three instructions (BPF\_SDIV, BPF\_SMOD, BPF\_MOVSX) have a non-zero offset.

The division and modulo operations support both unsigned and signed flavors.

For unsigned operations (BPF\_DIV and BPF\_MOD), for BPF\_ALU, 'imm' is interpreted as a 32-bit unsigned value. For BPF\_ALU64, 'imm' is first **:term:sign extended<Sign Extend>** from 32 to 64 bits, and then interpreted as a 64-bit unsigned value.

For signed operations (BPF\_SDIV and BPF\_SMOD), for BPF\_ALU, 'imm' is interpreted as a 32-bit signed value. For BPF\_ALU64, 'imm' is first **:term:sign extended<Sign Extend>** from 32 to 64 bits, and then interpreted as a 64-bit signed value.

Note that there are varying definitions of the signed modulo operation when the dividend or divisor are negative, where implementations often vary by language such that Python, Ruby, etc. differ from C, Go, Java, etc. This specification requires that signed modulo use truncated division (where  $-13 \% 3 == -1$ ) as implemented in C, Go, etc.:

$$a \% n = a - n * \text{trunc}(a / n)$$

The BPF\_MOVSX instruction does a move operation with sign extension. BPF\_ALU | BPF\_MOVSX **:term:sign extends<Sign Extend>** 8-bit and 16-bit operands into 32 bit operands, and zeroes the remaining upper 32 bits. BPF\_ALU64 | BPF\_MOVSX **:term:sign extends<Sign Extend>** 8-bit, 16-bit, and 32-bit operands into 64 bit operands.

Shift operations use a mask of 0x3F (63) for 64-bit operations and 0x1F (31) for 32-bit operations.

### 1.3.2 Byte swap instructions

The byte swap instructions use instruction classes of BPF\_ALU and BPF\_ALU64 and a 4-bit 'code' field of BPF\_END.

The byte swap instructions operate on the destination register only and do not use a separate source register or immediate value.

For BPF\_ALU, the 1-bit source operand field in the opcode is used to select what byte order the operation converts from or to. For BPF\_ALU64, the 1-bit source operand field in the opcode is reserved and must be set to 0.

class	source	value	description
BPF_ALU	BPF_TO_LE	0x00	convert between host byte order and little endian
BPF_ALU	BPF_TO_BE	0x08	convert between host byte order and big endian
BPF_ALU64	Reserved	0x00	do byte swap unconditionally

The 'imm' field encodes the width of the swap operations. The following widths are supported: 16, 32 and 64.

Examples:

BPF\_ALU | BPF\_TO\_LE | BPF\_END with imm = 16/32/64 means:

```
dst = htobe16(dst)
dst = htobe32(dst)
dst = htobe64(dst)
```

BPF\_ALU | BPF\_TO\_BE | BPF\_END with imm = 16/32/64 means:

```
dst = htole16(dst)
dst = htole32(dst)
dst = htole64(dst)
```

BPF\_ALU64 | BPF\_TO\_LE | BPF\_END with imm = 16/32/64 means:

```
dst = bswap16(dst)
dst = bswap32(dst)
dst = bswap64(dst)
```

### 1.3.3 Jump instructions

BPF\_JMP32 uses 32-bit wide operands while BPF\_JMP uses 64-bit wide operands for otherwise identical operations. The 'code' field encodes the operation as below:

code	value	src	description	notes
BPF_JA	0x00	0x00	PC += offset	BPF_JMP class
BPF_JA	0x00	0x00	PC += imm	BPF_JMP32 class
BPF_JEQ	0x01	any	PC += offset if dst == src	
BPF_JGT	0x02	any	PC += offset if dst > src	unsigned
BPF_JGE	0x03	any	PC += offset if dst >= src	unsigned
BPF_JSET	0x04	any	PC += offset if dst & src	
BPF_JNE	0x05	any	PC += offset if dst != src	
BPF_JSGT	0x06	any	PC += offset if dst > src	signed
BPF_JSGE	0x07	any	PC += offset if dst >= src	signed
BPF_CALL	0x08	0x00	call helper function by address	see <a href="#">Helper functions</a>
BPF_CALL	0x08	0x01	call PC += imm	see <a href="#">Program-local functions</a>

BPF_CALL	0x8	0x2	call helper function by BTF ID	see <a href="#">Helper functions</a>
BPF_EXIT	0x9	0x0	return	BPF_JMP only
BPF_JLT	0xa	any	PC += offset if dst < src	unsigned
BPF_JLE	0xb	any	PC += offset if dst <= src	unsigned
BPF_JSLT	0xc	any	PC += offset if dst < src	signed
BPF_JSLE	0xd	any	PC += offset if dst <= src	signed

The BPF program needs to store the return value into register R0 before doing a BPF\_EXIT.

Example:

BPF\_JSGE | BPF\_X | BPF\_JMP32 (0x7e) means:

```
if (s32)dst s>= (s32)src goto +offset
```

where 's>=' indicates a signed '>=' comparison.

BPF\_JA | BPF\_K | BPF\_JMP32 (0x06) means:

```
goto1 +imm
```

where 'imm' means the branch offset comes from insn 'imm' field.

Note that there are two flavors of BPF\_JA instructions. The BPF\_JMP class permits a 16-bit jump offset specified by the 'offset' field, whereas the BPF\_JMP32 class permits a 32-bit jump offset specified by the 'imm' field. A > 16-bit conditional jump may be converted to a < 16-bit conditional jump plus a 32-bit unconditional jump.

#### 1.3.3.1 Helper functions

Helper functions are a concept whereby BPF programs can call into a set of function calls exposed by the underlying platform.

Historically, each helper function was identified by an address encoded in the imm field. The available helper functions may differ for each program type, but address values are unique across all program types.

Platforms that support the BPF Type Format (BTF) support identifying a helper function by a BTF ID encoded in the imm field, where the BTF ID identifies the helper name and type.

#### 1.3.3.2 Program-local functions

Program-local functions are functions exposed by the same BPF program as the caller, and are referenced by offset from the call instruction, similar to BPF\_JA. The offset is encoded in the imm field of the call instruction. A BPF\_EXIT within the program-local function will return to the caller.



## 1.4 Load and store instructions

For load and store instructions (BPF\_LD, BPF\_LDX, BPF\_ST, and BPF\_STX), the 8-bit 'opcode' field is divided as:

3 bits (MSB)	2 bits	3 bits (LSB)
mode	size	instruction class

The mode modifier is one of:

mode modifier	value	description	reference
BPF_IMM	0x00	64-bit immediate instructions	<a href="#">64-bit immediate instructions</a>
BPF_ABS	0x20	legacy BPF packet access (absolute)	<a href="#">Legacy BPF Packet access instructions</a>
BPF_IND	0x40	legacy BPF packet access (indirect)	<a href="#">Legacy BPF Packet access instructions</a>
BPF_MEM	0x60	regular load and store operations	<a href="#">Regular load and store operations</a>
BPF_MEMSX	0x80	sign-extension load operations	<a href="#">Sign-extension load operations</a>
BPF_ATOMIC	0xc0	atomic operations	<a href="#">Atomic operations</a>

The size modifier is one of:

size modifier	value	description
BPF_W	0x00	word (4 bytes)
BPF_H	0x08	half word (2 bytes)
BPF_B	0x10	byte
BPF_DW	0x18	double word (8 bytes)

### 1.4.1 Regular load and store operations

The BPF\_MEM mode modifier is used to encode regular load and store instructions that transfer data between a register and memory.

BPF\_MEM | <size> | BPF\_STX means:

```
*(size *) (dst + offset) = src
```

BPF\_MEM | <size> | BPF\_ST means:

```
*(size *) (dst + offset) = imm32
```

BPF\_MEM | <size> | BPF\_LDX means:

```
dst = *(unsigned size *) (src + offset)
```

Where size is one of: BPF\_B, BPF\_H, BPF\_W, or BPF\_DW and 'unsigned size' is one of u8, u16, u32 or u64.

### 1.4.2 Sign-extension load operations

The `BPF_MEMSX` mode modifier is used to encode **term: `sign-extension`** `<Sign Extend>` load instructions that transfer data between a register and memory.

`BPF_MEMSX | <size> | BPF_LDX` means:

```
dst = *(signed size *) (src + offset)
```

Where size is one of: `BPF_B`, `BPF_H` or `BPF_W`, and 'signed size' is one of `s8`, `s16` or `s32`.

### 1.4.3 Atomic operations

Atomic operations are operations that operate on memory and can not be interrupted or corrupted by other access to the same memory region by other BPF programs or means outside of this specification.

All atomic operations supported by BPF are encoded as store operations that use the `BPF_ATOMIC` mode modifier as follows:

- `BPF_ATOMIC | BPF_W | BPF_STX` for 32-bit operations
- `BPF_ATOMIC | BPF_DW | BPF_STX` for 64-bit operations
- 8-bit and 16-bit wide atomic operations are not supported.

The 'imm' field is used to encode the actual atomic operation. Simple atomic operation use a subset of the values defined to encode arithmetic operations in the 'imm' field to encode the atomic operation:

imm	value	description
<code>BPF_ADD</code>	<code>0x00</code>	atomic add
<code>BPF_OR</code>	<code>0x40</code>	atomic or
<code>BPF_AND</code>	<code>0x50</code>	atomic and
<code>BPF_XOR</code>	<code>0xa0</code>	atomic xor

`BPF_ATOMIC | BPF_W | BPF_STX` with 'imm' = `BPF_ADD` means:

```
*(u32 *) (dst + offset) += src
```

`BPF_ATOMIC | BPF_DW | BPF_STX` with 'imm' = `BPF_ADD` means:

```
*(u64 *) (dst + offset) += src
```

In addition to the simple atomic operations, there also is a modifier and two complex atomic operations:

imm	value	description
<code>BPF_FETCH</code>	<code>0x01</code>	modifier: return old value
<code>BPF_XCHG</code>	<code>0xe0   BPF_FETCH</code>	atomic exchange
<code>BPF_CMPXCHG</code>	<code>0xf0   BPF_FETCH</code>	atomic compare and exchange

The `BPF_FETCH` modifier is optional for simple atomic operations, and always set for the complex atomic operations. If the `BPF_FETCH` flag is set, then the operation also overwrites `src` with the value that was in memory before it was modified.

The `BPF_XCHG` operation atomically exchanges `src` with the value addressed by `dst + offset`.

The `BPF_CMPXCHG` operation atomically compares the value addressed by `dst + offset` with `R0`. If they match, the value addressed by `dst + offset` is replaced with `src`. In either case, the value that was at `dst + offset` before the operation is zero-extended and loaded back to `R0`.

#### 1.4.4 64-bit immediate instructions

Instructions with the `BPF_IMM` 'mode' modifier use the wide instruction encoding defined in [Instruction encoding](#), and use the 'src' field of the basic instruction to hold an opcode subtype.

The following table defines a set of `BPF_IMM` | `BPF_DW` | `BPF_LD` instructions with opcode subtypes in the 'src' field, using new terms such as "map" defined further below:

opcode construction	opcode	src	pseudocode	imm type	dst type
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x0	<code>dst = imm64</code>	integer	integer
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x1	<code>dst = map_by_fd(imm)</code>	map fd	map
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x2	<code>dst = map_val(map_by_fd(imm)) + next_imm</code>	map fd	data pointer
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x3	<code>dst = var_addr(imm)</code>	variable id	data pointer
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x4	<code>dst = code_addr(imm)</code>	integer	code pointer
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x5	<code>dst = map_by_idx(imm)</code>	map index	map
<code>BPF_IMM</code>   <code>BPF_DW</code>   <code>BPF_LD</code>	0x18	0x6	<code>dst = map_val(map_by_idx(imm)) + next_imm</code>	map index	data pointer

where

- `map_by_fd(imm)` means to convert a 32-bit file descriptor into an address of a map (see [Maps](#))
- `map_by_idx(imm)` means to convert a 32-bit index into an address of a map
- `map_val(map)` gets the address of the first value in a given map
- `var_addr(imm)` gets the address of a platform variable (see [Platform Variables](#)) with a given id
- `code_addr(imm)` gets the address of the instruction at a specified relative offset in number of (64-bit) instructions
- the 'imm type' can be used by disassemblers for display
- the 'dst type' can be used for verification and JIT compilation purposes

##### 1.4.4.1 Maps

Maps are shared memory regions accessible by BPF programs on some platforms. A map can have various semantics as defined in a separate document, and may or may not have a single contiguous memory region, but the '`map_val(map)`' is currently only defined for maps that do have a single contiguous memory region.

Each map can have a file descriptor (fd) if supported by the platform, where '`map_by_fd(imm)`' means to get the map with the specified file descriptor. Each BPF program can also be defined to use a set of maps associated with the program at load time, and '`map_by_idx(imm)`' means to get the map with the given index in the set associated with the BPF program containing the instruction.

#### 1.4.4.2 Platform Variables

Platform variables are memory regions, identified by integer ids, exposed by the runtime and accessible by BPF programs on some platforms. The 'var\_addr(imm)' operation means to get the address of the memory region identified by the given id.

#### 1.4.5 Legacy BPF Packet access instructions

BPF previously introduced special instructions for access to packet data that were carried over from classic BPF. However, these instructions are deprecated and should no longer be used.